

Infer AI Engineer Task

Timeline: 2 days

Voice Agent Evals

Livekit (<https://github.com/pipecat-ai/pipecat>) is an open-source framework for building AI voice agents. Tau Bench (<https://taubench.com/#home>) is a benchmark for voice agents. Pick a domain on tau bench and evaluate gpt 4.1 or any other model of your choice (including STT/TTS) against some of those tasks (Check audio trajectories for visualizing some audio samples and tasks for the relevant policy and tool calls). Build evals for 3 behaviors where the base model is failing and then improve on those through prompt changes.

Submission Deliverables

1. Github repo with brief technical write-up on the PR explaining:
 - Your implementation approach
 - Approaches considered/ trade offs made
 - Any future improvements
2. Loom videos/ recordings where you demo and explain your write up within 5mins
3. Claude code/codex session transcripts